

CIFAR10_LSTM

October 25, 2019

```
[166]: import numpy as np
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import time

[167]: # set random seeds for reproducibility
torch.manual_seed(12)
torch.cuda.manual_seed(12)
np.random.seed(12)

[168]: # Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# If we are on a CUDA machine, then this should print a CUDA device, but we are
→not, so it will run on CPU:
print(f'Working on device={device}')
```

Working on device=cpu

```
[169]: # Hyper-parameters

# each CIFAR image is RGB 32x32, so it is an 3D array [3,32,32]
# we will flatten the image as vector dim=3*32*32

# we have 10 classes
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

num_classes = 10

num_epochs = 10
batch_size = 1024

learning_rate = 0.001

[170]: transform = transforms.Compose(
    [transforms.ToTensor()],
```

```

        transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                       download=True, transform=transforms.
                                       ↳ToTensor())

train_loader = torch.utils.data.DataLoader(trainset, batch_size=1024,
                                       shuffle=True)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transforms.
                                       ↳ToTensor())

test_loader = torch.utils.data.DataLoader(testset,
                                       batch_size=1024,
                                       shuffle=False)

```

Files already downloaded and verified

Files already downloaded and verified

```

[171]: ies = []
       image = []
       label = []

       for i, (images, labels) in enumerate(train_loader):
           ies.append(i)
           image.append(images)
           label.append(labels)

```

```

[172]: image[0][0].view(96,32).size()

```

```

[172]: torch.Size([96, 32])

```

```

[173]: class LSTMModel(nn.Module):
       def __init__(self, input_dim, hidden_dim, layer_dim, output_dim):
           super(LSTMModel, self).__init__()
           # Hidden dimensions
           self.hidden_dim = hidden_dim

           # Number of hidden layers
           self.layer_dim = layer_dim

           # Building your LSTM
           # batch_first=True causes input/output tensors to be of shape
           # (batch_dim, seq_dim, feature_dim)
           self.lstm = nn.LSTM(input_dim, hidden_dim, layer_dim, batch_first=True)

           # Readout layer
           self.fc = nn.Linear(hidden_dim, output_dim)

```

```

def forward(self, x):
    # Initialize hidden state with zeros
    h0 = torch.zeros(self.layer_dim, x.size(0), self.hidden_dim).
    ↳requires_grad_()

    # Initialize cell state
    c0 = torch.zeros(self.layer_dim, x.size(0), self.hidden_dim).
    ↳requires_grad_()

    # 28 time steps
    # We need to detach as we are doing truncated backpropagation through
    ↳time (BPTT)
    # If we don't, we'll backprop all the way to the start even after going
    ↳through another batch
    out, (hn, cn) = self.lstm(x, (h0.detach(), c0.detach()))

    # Index hidden state of last time step
    # out.size() --> 100, 28, 100
    # out[:, -1, :] --> 100, 100 --> just want last time step hidden states!
    ↳
    out = self.fc(out[:, -1, :])
    # out.size() --> 100, 10
    return out

```

```

[174]: input_dim = 32
       hidden_dim = 128
       layer_dim = 1
       output_dim = 10

```

```

[175]: model = LSTMModel(input_dim, hidden_dim, layer_dim, output_dim)

```

```

[176]: # Loss and optimizer
       criterion = nn.CrossEntropyLoss()
       optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)

```

```

[177]: # Number of steps to unroll
       seq_dim = 96
       iter = 0

       for epoch in range(num_epochs):
           t0 = time.time()
           for i, (images, labels) in enumerate(train_loader):
               # Load images as a torch tensor with gradient accumulation abilities
               epoch_loss = 0
               images = images.view(-1, seq_dim, input_dim).requires_grad_()
               labels = labels.to(device)

```

```

# Clear gradients w.r.t. parameters
optimizer.zero_grad()

# Forward pass to get output/logits
# outputs.size() --> 100, 10
outputs = model(images)

# Calculate Loss: softmax --> cross entropy loss
loss = criterion(outputs, labels)
epoch_loss += loss.item()

# Getting gradients w.r.t. parameters
loss.backward()

# Updating parameters
optimizer.step()

iter += 1

if iter % 49 == 0:
    # Calculate Accuracy
    correct = 0
    total = 0
    # Iterate through test dataset
    for images, labels in test_loader:
        # Resize image
        images = images.view(-1, seq_dim, input_dim)

        # Forward pass only to get logits/output
        outputs = model(images)

        # Get predictions from the maximum value
        _, predicted = torch.max(outputs.data, 1)

        # Total number of labels
        total += labels.size(0)

        # Total correct predictions
        correct += (predicted == labels).sum().item()

    accuracy = correct / total

    # Print Loss
    print(

```

```
f'Epoch [{epoch+1}/{num_epochs}]], Iter: {iter} Loss: {loss.item():.
→4f} Test acc: {accuracy: .4f}')
```

```
print('{ seconds'.format(time.time() - t0))
```

```
Epoch [1/10]], Iter: 49 Loss: 2.0424 Test acc: 0.2571
90.77259087562561 seconds
Epoch [2/10]], Iter: 98 Loss: 1.8783 Test acc: 0.3171
103.57928109169006 seconds
Epoch [3/10]], Iter: 147 Loss: 1.8695 Test acc: 0.3493
101.47186303138733 seconds
Epoch [4/10]], Iter: 196 Loss: 1.7231 Test acc: 0.3665
99.46165418624878 seconds
Epoch [5/10]], Iter: 245 Loss: 1.7573 Test acc: 0.3736
100.66864204406738 seconds
Epoch [6/10]], Iter: 294 Loss: 1.6618 Test acc: 0.3871
101.9321620464325 seconds
Epoch [7/10]], Iter: 343 Loss: 1.6320 Test acc: 0.3965
99.57273006439209 seconds
Epoch [8/10]], Iter: 392 Loss: 1.6361 Test acc: 0.4086
96.93101382255554 seconds
Epoch [9/10]], Iter: 441 Loss: 1.6270 Test acc: 0.4058
100.76327204704285 seconds
Epoch [10/10]], Iter: 490 Loss: 1.6105 Test acc: 0.4018
97.01067399978638 seconds
```